


Building Local LLMs with Low-bit Quantization

Reducing Model Size Without Sacrificing Performance

PD Dr. Haojin Yang,
Multimedia and Machine Learning Group,
Hasso Plattner Institute

**Design IT.
Create Knowledge.**

www.hpi.de



Introduction to Low-Bit Quantization



- Low-bit quantization reduces model parameter precision to make LLMs smaller
- Converts 32-bit floating-point numbers to 8-bit, 4-bit, or even lower formats
- Enables running powerful LLMs on consumer hardware with limited resources
- Significantly reduces memory requirements and computational demands
- Trades minimal accuracy loss for substantial efficiency gains
- Essential technique for local deployment of AI models

How Quantization Works

- Quantization reduces precision of model weights from 32-bit floating point to lower bit formats (8-bit, 4-bit, etc.)
- Process maps continuous values to a smaller, discrete set of values using scaling factors (α, β)
- Post-training quantization (PTQ) applies quantization after model is fully trained
- Quantization-aware training (QAT) incorporates quantization during the training process
- **Weight-only quantization** focuses on model parameters while preserving activation precision
- Different methods balance trade-offs between model size, inference speed, and accuracy





Image source:
<https://www.maartengrootendorst.com/blog/quantization/>

Quantization Techniques Compared

- 8-bit quantization: $\sim 2\times$ smaller models with minimal performance loss
- 4-bit quantization: $\sim 8\times$ smaller models with acceptable quality tradeoffs
- 2-bit/3-bit quantization: Experimental with significant quality degradation
- GPTQ vs AWQ: Popular algorithms with different optimization approaches
- INT4 vs NF4: Different number formats with varying performance characteristics
- Mixed-precision: Using different bit-widths for different model components
 - Layer-level mixed bit-width or channel-level mixed bit-width

Model size estimation of a 70B Model with 70 billion parameters:

64-bit = $64/8 \times 70\text{B} \approx 560 \text{ GB}$

32-bit = $32/8 \times 70\text{B} \approx 280 \text{ GB}$

16-bit = $16/8 \times 70\text{B} \approx 140 \text{ GB}$

8-bit = $8/8 \times 70\text{B} \approx 70 \text{ GB}$

4-bit = $4/8 \times 70\text{B} \approx 35 \text{ GB}$

2-bit = $2/8 \times 70\text{B} \approx 17,7 \text{ GB}$

Symmetric Quantization




- Maps floating-point values to a symmetric range around zero
- Zero in floating-point space maps exactly to zero in quantized space
- Uses a linear mapping centered around zero

- Scale factor: $S = \frac{(2^{b-1} - 1)}{\alpha}$ (where α is the maximum absolute value)

- Quantization: $q(x) = \text{round}(s \times x)$

- Dequantization: $\hat{x} = \frac{q(x)}{s}$



Symmetric Quantization Example


- Example with original values: [8.5, 3.08, 3.02, -7.59, 10.8]
- Highest absolute value: 10.8
- Scale factor calculation: $S = 127 / 10.8 = 11.76$
- Quantized values: [100, 36, 36, -89, 127]
- **Advantages:**
 - Simple implementation
 - Preserves zero exactly
 - No need to store zero-point offset
- **Limitations:**
 - Less efficient for data not centered around zero
 - Similar values may map to same quantized value (loss of precision)

Asymmetric Quantization



- Not symmetric around zero
- Maps minimum (β) and maximum (α) values to full quantized range
- Uses zero-point quantization with shifted zero position

- Scale factor: $s = \frac{(2^b - 1)}{(\alpha - \beta)}$
- Zero point: $z = \text{round}(-\beta \times s)$
- Quantization: $q = \text{round}(x \times s + z)$
- Dequantization: $x = \frac{(q - z)}{s}$



Comparing Symmetric vs. Asymmetric Quantization

- Symmetric: Centered around zero, simpler calculations
- Asymmetric: Shifted zero position, better range utilization
- When to use Symmetric:
 - Data naturally centered around zero (like weights in neural networks)
 - When computational simplicity is important
 - When preserving zero exactly is critical
- When to use Asymmetric:
 - Data with skewed distribution (not centered around zero)
 - When maximizing representation accuracy is priority
 - For activations in neural networks (often positive-only)